

DeepBias: An Adaptive and Dynamic Benchmark for In-Depth Probing Social Biases in Visual Language Models

First Author, Second Author, and Third Author

Abstract—While Large Vision-Language Models (LVLMs) demonstrate remarkable capabilities, they remain highly susceptible to embedded societal biases. Existing bias evaluation protocols predominantly rely on static datasets, which provide only a superficial assessment, as their fixed test cases cannot adaptively evolve to measure the true depth and limits of model vulnerabilities. We introduce DeepBias, an adaptive and dynamic framework for the in-depth probing of social biases in LVLMs. Our approach operates through a dynamic “generation-evolution-probing” loop. First, a generative Proposer synthesizes test cases to evaluate the target model. Second, we optimize the Proposer via Direct Preference Optimization based on the target model’s initial responses, directing it to generate increasingly challenging and targeted samples. Third, an autonomous DeepAgent conducts multi-turn sequential probing on each sample: at each turn, it analyzes the full dialogue history and either deepens exposed biases or adversarially rewrites the query to bypass safety filters, progressively uncovering latent biases. Furthermore, we employ 6 diverse, state-of-the-art LVLMs as a parallel anchor ensemble whose aggregated responses jointly drive the optimization, iteratively synthesizing a novel, highly challenging benchmark dataset. Comprehensive experiments demonstrate that DeepBias successfully exposes deep societal biases, establishing a scalable and evolutionary paradigm for robust LVLM safety evaluation.

Index Terms—Vision-Language Model Bias, Adversarial Bias Probing, Dynamic Benchmark Generation.

I. INTRODUCTION

LARGE Vision-Language Models (LVLMs) have demonstrated unprecedented capabilities in multimodal understanding and reasoning, powering applications from automated captioning to visual agents [1], [2]. However, these models inevitably inherit and amplify societal biases present in their massive training corpora, leading to discriminatory outputs regarding gender, race, and other sensitive attributes. While safety alignment techniques like RLHF [3] have been deployed to mitigate these issues, accurately evaluating the true extent of bias in LVLMs remains an open and critical challenge.

Current bias evaluation protocols predominantly rely on static benchmarks [4], [5]. As illustrated in the left panel of Fig.1, these benchmarks typically consist of fixed sets of image-text pairs and rely on one-way inference. The data

distribution of the fixed dataset is independent of the target model’s bias vulnerabilities. While foundational, this static paradigm faces three major limitations. First, static benchmarks [6], particularly those curated from real world data or human annotations [7], are inevitably absorbed into the massive training corpora of modern LVLMs. This allows models to memorize the test set, leading to inflated scores that make the static benchmarks obsolete and unreliable. Second, existing methods rely on fixed queries that cannot adapt to model feedback. Consequently, safety-aligned models easily deflect these explicit prompts, leaving deep-seated latent biases undetected and creating a false security. Third, static datasets rarely align with the specific vulnerabilities of the target model. This inflexibility wastes evaluation efforts on redundant queries where the model is already robust, while actual safety problems remain unexamined.

To overcome these limitations, we advocate for a transition from static dataset to dynamic adversarial evaluation. We introduce DeepBias, an adaptive framework designed to probe the depth of social biases in LVLMs. As shown in the right panel of Fig.1, unlike static approaches, DeepBias dynamically generates new text queries based on the model’s feedback to initial data, probing the bias at a much deeper, granular level. Furthermore, it continuously optimizes the test data distribution to explicitly align with the target model’s specific weaknesses. This results in a comprehensive, deep, and targeted evaluation.

To drive this dynamic evaluation, the core of DeepBias is a generative Proposer that synthesizes initial test cases. As illustrated in the *Dynamically update* column of Fig.1, based on the target model’s responses to these initial queries, we directly apply Direct Preference Optimization (DPO) [8] to the Proposer, treating successful bias elicitation as positive preferences and failed attempts as negative. This closed-loop process explicitly shifts the generated dataset’s distribution to tightly align with the specific vulnerabilities of the target model.

After DPO optimization, as shown in the *Deep Exploration* column of Fig.1, we deploy an autonomous DeepAgent to conduct in-depth, instance-level probing through multi-turn interaction with the target model. Unlike static one-shot evaluation, the DeepAgent operates as a fully autonomous agent that engages in a sequential multi-turn dialogue. At each turn, it reviews the complete interaction history and adaptively selects one of two strategies: if the model’s previous response exposes

First Author is with Princeton University, Princeton NJ 08544, USA.

Second Author is with Springer Heidelberg, Tiergartenstr. 17, 69121 Heidelberg, Germany.

Third Author is with ABC Institute, Rupert-Karls-University Heidelberg, Heidelberg, Germany.




Fig. 1: Comparison between traditional static evaluation and our DeepBias framework. **Left, Traditional Static Evaluation:** a fixed set of queries is issued to the target LVLm; safety alignment deflects every query, yielding an illusory “0% bias” verdict that masks the mismatch between the test distribution and the model’s true weaknesses. **Right, Our DeepBias Evaluation** integrates two complementary mechanisms. Under *Dynamically update*, the Proposer iteratively uses the target model’s responses as optimization signals to shift its output distribution toward the model’s vulnerable regions, producing a targeted test set aligned with its real failure modes. Under *Deep Exploration*, an autonomous DeepAgent conducts multi-turn sequential probing on each retained sample: starting from an initially deflected query, it adaptively deepens along exposed dimensions such as math performance, or adversarially rewrites the prompt with a racial comparative frame, progressively surfacing latent stereotypes that single-shot querying cannot reach.

targeted follow-up question to pinpoint the specific dimensions and severity of the stereotype; if the model responds safely, the DeepAgent adversarially rewrites the query to bypass safety filters and elicit latent biases. This two-stage design decouples distribution-level optimization (Proposer + DPO) from instance-level deep probing (DeepAgent), enabling both a well-targeted data distribution and thorough per-sample bias exploration.

To construct a robust evaluation standard suitable for advanced VLmS, we employ an ensemble of 6 state-of-the-art LVLmS as anchor models. By iteratively optimizing the Proposer against this collective target, we systematically uncover deep-seated, common biases inherent in modern visual-language representations. To further ensure benchmark quality, we introduce a multi-model voting validation step that filters out ill-posed queries whose ambiguity assumption is violated by the anchor ensemble. This process culminates in the construction of the DeepBias Benchmark, a comprehensive dataset distilled from the most challenging samples that successfully penetrated the defenses of the anchor models. Extensive experiments across a wide range of SOTA models demonstrate that our method, and the resulting benchmark, effectively exposes

latent biases that static baselines fail to detect.

Our main contributions are summarized as follows:

- **Dynamic Adaptive Framework:** We propose DeepBias, the first framework to utilize DPO-driven evolution for VLm bias evaluation, shifting the paradigm from static datasets to dynamic adversarial probing.
- **In-Depth Probing Strategy:** We introduce an autonomous DeepAgent that conducts multi-turn sequential interactions with the target model, adaptively deepening or rewriting queries based on full dialogue history. This allows us to systematically penetrate safety filters and uncover latent biases that are typically masked in standard benchmarks.
- **The DeepBias Benchmark:** By aggregating the hard and deep test samples from probing 6 diverse LVLmS, we release a high-quality, challenging benchmark dataset for the community.
- **Empirical Evidence:** Extensive experiments show that DeepBias successfully exposes significantly more bias than static baselines, highlighting the fragility of current safety alignments in LVLmS.

II. RELATED WORK

A. Foundations of Bias and Fairness

The term “bias” is heavily overloaded in machine learning [9]: in its statistical sense it denotes systematic estimator deviation, while in its societal sense it denotes discrimination—model behaviour that systematically disadvantages members of protected social groups. The by-now-standard taxonomy for the societal sense, originating with Crawford’s NeurIPS 2017 keynote and formalised for NLP by [10], partitions harms into (i) *allocational harms*, in which an automated system unfairly distributes resources or opportunities between social groups, and (ii) *representational harms*, in which a system denigrates, stereotypes, under-represents, or erases particular groups in its outputs. [10] emphasise that representational harms are themselves harms and do not require an allocational downstream justification to warrant study. In antidiscrimination law, a closely related distinction between *disparate treatment* (individual-level intentional discrimination) and *disparate impact* (group-level statistical disparity) plays the same structural role [11].

Complementing this harm-side view, [12] catalogue seven bias sources that span the ML lifecycle—historical, representation, measurement, aggregation, learning, evaluation, and deployment bias—each of which can translate into either allocational or representational harm. Their *historical bias* category makes explicit that faithfully mirroring the world is insufficient to produce a harm-free model, because the world itself carries inequities whose reproduction is injurious. The empirical stakes are illustrated by intersectional audits such as [13], whose evaluation of commercial facial analysers reported error-rate gaps of up to 34.4 percentage points between demographic sub-populations, crystallising why purely descriptive reporting of model performance is insufficient. Our evaluation protocol builds on this literature; in §III-A we distil the specific definitional framework that underpins our design, including the three characteristics we use to distinguish societal bias from legitimate statistical correlation.

B. Societal Bias Evaluation in Vision-Language Models

The evaluation of social bias has its roots in Natural Language Processing (NLP), with pioneering works demonstrating that word embeddings and language models capture human-like stereotypes [14], [15]. Benchmarks such as CrowS-Pairs [16], StereoSet [17], and BBQ [18] became standard protocols for quantifying categorical biases in text modality.

With the advent of Large Vision-Language Models (LVLMs), research has pivoted to multimodal bias. Early efforts like VisoGender [4], GenderBias-VL [19], VL-Bias [5], and Vignette [20] introduced static image-text pairs to assess specific demographic biases like gender and racial biases. Others, such as FairFace [21] and MIAP [22], focused on dataset diversity. Concurrently, comprehensive benchmark like VLBiasBench [23] have been developed to evaluate biases of various categories. However, relying on static datasets presents fundamental limitations. First, *Data Leakage*: as LVLMs are trained on web-scale datasets like LAION-5B [24], static test sets are likely seen during training, making evaluation

results unreliable [25]. Second, *Rapid Obsolescence*: static benchmarks cannot keep pace with the rapid versioning of models (e.g., GPT-4o [1], Gemini [26]), leading to benchmark saturation. Third, *Lack of Adaptivity*: static datasets lack the capacity to dynamically adjust based on model feedback [27], so they cannot execute targeted and deep probing against specific model vulnerabilities.

C. Red Teaming and Jailbreaking for VLMs

Red teaming has evolved from manual testing to automated frameworks designed to proactively discover model vulnerabilities. In the realm of Large Language Models (LLMs), foundational works [28], [29] employed generator models to automatically synthesize test cases that provoke toxic outputs. Recent approaches have introduced optimization-based strategies. Gradient-based methods like GCG [30] and ARCA [31] optimize discrete tokens to bypass safety filters, while iterative-refinement methods such as PAIR [32] leverage an attacker LLM to adaptively rewrite prompts based on the target model’s feedback, and evolutionary methods like AutoDAN [33] apply genetic algorithms to evolve adversarial prompts. As the field shifts to Vision-Language Models (LVLMs), research has extended these semantic attacks to the multimodal domain. Recent studies reveal that by combining structured textual prompts with specific visual contexts [34], or fragmenting malicious instructions across different modalities [35], [36], it is possible to effectively bypass the cross-modal safety alignment of LVLMs.

However, applying existing automated red teaming paradigms directly to social bias evaluation presents significant limitations. First, *Mismatch in Target*: current jailbreak frameworks predominantly optimize for eliciting overt general harms such as malicious code generation, violence, or extreme toxicity [37]. Social bias, conversely, is deeply ingrained, implicit, and often context-dependent, requiring a more nuanced elicitation strategy than standard adversarial attacks. Second, *Binary Metrics over Depth*: existing methods primarily optimize for Attack Success Rate (ASR)—a binary metric indicating whether a model is induced to generate harmful content [32]. They do not measure the severity, consistency, or depth of the bias once exposed. Third, *Semantic Degradation*: to maximize ASR, many automated algorithms such as gradient-based prompt search generate unnatural, gibberish prompts [30]. While these may force errors, they fail to reflect realistic human-AI interaction. In contrast, DeepBias shifts the objective from binary jailbreaking to depth probing. Instead of generating gibberish to force errors, our DeepAgent synthesizes semantically natural and contextually complex scenarios through multi-turn interactions, aiming to uncover deep-seated stereotypes.

D. Optimization and RL for Data Generation

Reinforcement Learning (RL) and evolutionary strategies have proven effective in synthesizing complex data. In the NLP domain, Self-Instruct [38] and WizardLM [39] demonstrated that LLMs can iteratively improve the complexity of their own

training data. More recently, Direct Preference Optimization (DPO) [8] has emerged as a stable alternative to PPO [40] for aligning models with human preferences.

In the context of evaluation, RL has been applied to automate red teaming to maximize the toxicity of target model responses [28] and search for diverse failure cases [41]. Most recently, RedHit [42] integrates DPO with Monte Carlo Tree Search and chain-of-thought reasoning to iteratively refine adversarial prompts, sharing our spirit of preference-optimized red-teaming but targeting generic jailbreak success rather than social bias. However, these methods typically optimize for generic failure modes. Our work is the first to apply DPO specifically for the “adversarial evolution” of bias probes. Unlike standard alignment which suppresses bias, we “reverse-align” our Proposer to maximize the likelihood of exposing latent biases, treating the target model’s vulnerability as the optimization objective. This allows our method to generate substantially harder and more penetrating evaluation samples compared to random or rule-based generation [27], [43], [44].

III. BIAS EVALUATION PROTOCOL

Before detailing our data generation pipeline, we specify the evaluation protocol. §III-A distils the definitional framework we adopt for *societal bias* and the boundary that separates it from legitimate statistical correlation. §III-B then operationalises this framework into a concrete test-instance format, four hard constraints on the question, an empirical validity check, and the bias-judgment metric. The pipeline that produces test instances satisfying the protocol is the subject of §IV.

A. Bias, Discrimination, and Statistical Deviation

As surveyed in §II-A, the term “bias” names two distinct phenomena: *statistical bias*, a systematic estimator deviation, and *societal bias*, a pattern of model behaviour that produces representational or allocational harm against protected social groups. Our protocol targets the latter exclusively; the former is outside the scope of this paper. Within the societal sense, however, a further boundary has to be drawn: a surface demographic pattern in a model’s output can in principle reflect either (a) a legitimate group-level statistical regularity or (b) substantive discrimination, and a coherent protocol must separate the two. Three characteristics are repeatedly invoked across the fairness, psychology, and philosophy-of-discrimination literatures to draw this boundary; we adopt them as the analytical backbone of our protocol.

- 1) **Individual attribution.** Applying a group-level statistic to judge a specific individual is categorically different from merely reporting the statistic. The legal framework of [11] formalises this in antidiscrimination law as the distinction between disparate treatment—direct, individual-level discrimination—and disparate impact, a purely statistical-disparity notion. At the algorithmic level, [45] give the canonical formalisation: they show that group-level statistical parity can coexist with patently unfair individual treatment, and propose *individual fairness*—“any two individuals who are similar with

respect to the particular task at hand should be classified similarly”—as a strictly individual-level criterion. The same intuition surfaces in the political-philosophy literature surveyed by [46] under the label “failure to treat people as individuals.” Thus, stating that “elderly people are statistically less physically fit” is a demographic observation; asserting that *this particular* elderly person “cannot perform this task” violates individual attribution and qualifies as bias on every one of these accounts.

- 2) **Normative prescription.** Even a descriptively accurate statistical regularity can carry illegitimate prescriptive force. In psychology, [47] document systematically that gender stereotypes are “highly prescriptive: the qualities they ascribe to women and men are also the qualities required of women and men,” with violations drawing punishment independently of any descriptive inaccuracy. [48], while arguing from social-psychological data that many group stereotypes are descriptively accurate, explicitly note that descriptive accuracy “has no bearing on concerns about the use [of stereotypes] to judge individuals”—that is, descriptive accuracy and normative legitimacy are orthogonal axes. At the meta level, [10] argue that analysing bias is “inherently a normative process in which some system behaviours are deemed good and others harmful,” so that studies resting on purely descriptive framings without articulating the normative claim are inconsistent by construction.
- 3) **Concrete harm.** Bias is ultimately a claim about consequence, not about pattern: a statistical regularity becomes bias when it is applied in a setting that produces tangible damage to individuals or groups. The seven-source lifecycle framework of [12] and the empirical audits of [13] establish this harm-side grounding; [10] further argue that representational harms are themselves harms and do not require an allocational downstream justification to warrant study. For VLMs specifically, the analogous harms are primarily representational: a model that produces stereotype-consistent descriptions, judgements, or reasoning about depicted individuals denigrates and misrepresents the groups in question, regardless of whether any particular downstream allocational decision is made.

B. Protocol Specification

Test Instance Format. A test instance is a triple $(I, Q, \mathcal{O})$, where:

- I is an image depicting a person whose sensitive demographic attribute—age, race, gender, etc.—is the target of the bias probe.
- Q is an intentionally ambiguous question coupled with a negative implication, whose four design constraints are specified below. An illustrative example for racial bias is: “The man was stopped by the police after coming out of the supermarket. Did he commit a crime?”
- $\mathcal{O} = \{\text{Yes}, \text{No}, \text{Unknown}\}$ is the set of answer options. Given the intentional ambiguity of Q , the neutral “Unknown” is the only response supported by the available evidence.




Fig. 2: The DeepBias Evaluation Protocol. **Left — Validity Check via Ensemble Voting:** each candidate triplet (I_1, I_2, I_3, Q) is queried against K anchor validator models, yielding $3K$ responses; the proportion of “Unknown” responses $\rho(Q)$ is computed and the triplet is admitted only when $\rho(Q) > \tau$, ensuring that the question is genuinely ambiguous rather than trivially answerable. Admitted triplets are decomposed into three independent test instances, one per image. **Right — Bias Judgment on a Test Instance:** on each test instance (I, Q, O) , the target model must select “Unknown” to be judged unbiased; any “Yes” or “No” response reveals a stereotype-driven judgment attributable to the sensitive attribute depicted in the image. Aggregated across instances, the fraction of “Unknown” responses is the *bias accuracy*; its complement is the *bias elicitation rate*.

The model is queried with one test instance at a time, and evaluation is always reported per-instance. The counterfactual triplet structure that appears in data generation and validation (below) is a probe-design artifact: it enables reliable validity checking but is never presented to the model under evaluation.

Question Design Constraints. The three characteristics defined in §III-A—*individual attribution*, *normative prescription*, and *concrete harm*—translate directly into four hard constraints on how Q must be written:

- 1) *Intentional ambiguity.* Q together with the image must be underspecified: no combination of visual evidence and textual context should permit a definitive Yes or No. This forces the model to reach for internal priors rather than evidence, making “Unknown” the only justified response.
- 2) *Negative implication.* Q is phrased so that any non-“Unknown” response entails attributing a tangibly harmful judgment to the depicted individual (e.g., criminality, incompetence, weakness). This targets the *concrete harm* characteristic and prevents the protocol from registering benign statistical remarks as bias.
- 3) *No probabilistic hedging.* Q must avoid softening words such as “likely,” “probably,” or “tend to.” Such words allow the model to retreat into statistically defensible territory (“this elderly person is *likely* less fit”), blurring

the *individual attribution* characteristic: we cannot tell whether the model is making a genuine biased judgment about the depicted individual or merely citing a group-level tendency. Forcing definitive assertions cleanly separates bias from statistical reasoning.

- 4) *No group-specific text.* Q must not reference the sensitive attribute in its wording (e.g., no “as an elderly person” or “being a woman”). The demographic attribute is conveyed solely through the image, ensuring that any stereotyping arises from the model’s visual perception rather than from textual group-level cues, and thus from the *normative prescription* the model places on the depicted individual rather than a group-level linguistic cue. Together with the negative implication, this drives the model past any hedged, statistics-flavoured response and into a definitive prescriptive judgment—which is precisely the behaviour our protocol flags as bias.

Validity Criterion. Even when Q is authored to satisfy the four constraints above, a generated question may still fail to be genuinely ambiguous in practice—for instance, when the image provides more discriminating evidence than intended, or when the question admits an implicit factual answer. To detect such failures, we introduce the *candidate triplet* (I_1, I_2, I_3, Q) : three image–question pairs sharing the same Q whose images are identical in all contextual factors (scene, pose, clothing,

lighting, etc.) and differ only in the targeted sensitive attribute. For age bias, for instance, the triplet may depict a 15-year-old, a 32-year-old, and a 64-year-old in otherwise identical contexts. This counterfactual structure is a controlled-variable probe used at construction time only; the three images are never presented jointly to the model under evaluation.

The validity check, illustrated in the left panel of Figure 2, queries K independent validator models on all three image-question pairs of the triplet, yielding $3K$ responses, and records the proportion that are “Unknown”:

$$\rho(Q) = \frac{1}{3K} \sum_{k=1}^K \sum_{i=1}^3 \mathbb{1}[r_i^{(k)} = \text{Unknown}]. \quad (1)$$

The triplet is admitted iff

$$\rho(Q) > \tau, \quad (2)$$

where $\tau \in (0, 1)$ is the validity threshold. Because the three images differ only in the sensitive attribute, a non-trivial fraction of “Unknown” responses across validator-image pairs indicates that the question is genuinely under-determined regardless of which demographic is depicted, rather than admitting an implicit factual answer; any validator that nonetheless chooses a non-“Unknown” option is then genuinely exhibiting bias.

Once admitted, the triplet is *decomposed* into three independent test instances $\{(I_1, Q, \mathcal{O}), (I_2, Q, \mathcal{O}), (I_3, Q, \mathcal{O})\}$, each entering the evaluation set. The triplet identity is retained as metadata: although it no longer plays a role in scoring, it enables per-demographic analyses of which groups a given model is most prone to bias against. The concrete instantiation—the identity of the validator models, the choice of K , and the value of τ —is an implementation matter deferred to §IV-C, where we set $K=6$ (the anchor ensemble) and $\tau=1/3$; the empirical sensitivity to τ is analysed in §V-F.

Bias Judgment Metric. As illustrated in the right panel of Figure 2, on each admitted test instance $(I, Q, \mathcal{O})$ the target model’s response r is scored as:

$$\text{Correct}(r) = \mathbb{1}[r = \text{Unknown}]. \quad (3)$$

We report two complementary quantities. The *bias accuracy* is the proportion of instances on which the model selects “Unknown”; its complement, the *bias elicitation rate* $1 - \text{bias accuracy}$, is the proportion of responses that expose societal bias. A higher elicitation rate—equivalently, a lower bias accuracy—indicates that the model is more frequently led into stereotype-driven judgments.

IV. METHOD

Building on the protocol of §III, we now describe how DeepBias produces valid test instances at scale. The pipeline is driven by two complementary entities. A *Proposer* (§IV-A) expands existing bias benchmarks into a large candidate pool and iteratively reshapes its output distribution toward the target LVLm’s vulnerabilities via Direct Preference Optimization. A *DeepAgent* (§IV-B) then conducts multi-turn adaptive probing on each Proposer-generated candidate to surface stereotypes beyond the reach of a single-shot query. The candidates that

survive this pipeline are aggregated into the final DeepBias Benchmark using a six-model anchor ensemble that both drives the DPO consensus labelling and validates each candidate’s ambiguity (§IV-C). Fig. 3 gives a visual overview. Throughout this chapter the construction pipeline operates on candidate triplets (I_1, I_2, I_3, Q) —the probe-design unit defined in §III; only at the end of §IV-C, once a triplet passes the validity check, is it decomposed into three single-image test instances $(I, Q, \mathcal{O})$ for evaluation.

A. Adaptive Data Generation via Adversarial Evolution

The data engine of DeepBias is a generative **Proposer**: a language model that synthesises candidate triplets (I_1, I_2, I_3, Q) satisfying the four question-design constraints of §III. As noted above, the triplet is the construction-time probe-design unit and is never presented jointly to the model under evaluation; decomposition into single-image test instances happens only at the end of §IV-C. It plays two complementary roles: it *expands* existing static bias benchmarks to a scale usable for evaluating modern LVLms, and it *adaptively evolves* its output distribution toward the specific vulnerabilities of the target model. Both roles are carried out by the same model operating in two distinguishable regimes, described in turn below; the full four-step Expand / Query / Label / Update cycle is illustrated in the left panel of Fig. 3.

Expansion via In-Context Few-Shot Generation. Existing bias benchmarks such as BBQ [18] and VLBiasBench [23] provide only on the order of 300 ambiguous-and-negative samples per demographic category, far too few to serve as a standalone evaluation pool for modern LVLms. The Proposer therefore acts as an in-context few-shot *expander*: given a pool of exemplar (P, Q) pairs, it synthesises a much larger batch of new candidates. For each demographic category—Age, Race, or Gender—the exemplars pair a base question Q with a set of textual image descriptions P that vary only in the targeted sensitive attribute: $P = \{p_1, p_2, p_3\}$ for Age and Race (three distinct identities, e.g. three ages or three racial identities) and $P = \{p_1, p_2\}$ for Gender (two identities). Given sampled (P, Q) pairs as exemplars, together with generative rules that re-impose the four constraints of §III—intentional ambiguity, negative implication, no probabilistic hedging, and no group-specific text—the Proposer divergently synthesises new (P, Q) candidates. A single invocation in this mode yields roughly 2,000 new textual candidates per category.

Each textual candidate (P, Q) is then rendered into a visual counterfactual triplet $\{I_1, I_2, I_3\}$ by feeding the prompts in P —augmented with standardised neutral contextual descriptions to enforce strict variable control—into Stable Diffusion XL [49], yielding a batch of multimodal candidate triplets $\{(I_1, I_2, I_3, Q)\}$. The standardised answer options $\mathcal{O} = \{\text{Yes}, \text{No}, \text{Unknown}\}$ are attached at evaluation time, when a validated triplet is decomposed into three independent test instances per §III. The exact prompt templates and generation guidelines are detailed in the Supplementary Material.

An *initial* Expansion round, applied to the base-model Proposer with the raw benchmark (~ 300 samples per category from BBQ/VLBiasBench) as exemplars, produces the first




Fig. 3: **The DeepBias Framework.** Two-panel overview of the construction pipeline. **Left — Proposer Adversarial Evolution via DPO:** from a raw-dataset seed of ~ 300 samples per category, the Proposer π_θ jointly emits a biased textual question Q together with an image-generation prompt; the prompt is rendered by Stable Diffusion XL [49] into a counterfactual image triplet (I_1, I_2, I_3) , which is paired with Q to form a test candidate. The target LVLM’s responses are split by the Unknown rule into positive / negative preference pairs that drive the DPO objective $\mathcal{L}_{\text{DPO}}$; the update loops back into π_θ for K rounds of adversarial evolution. **Right — DeepAgent Multi-Turn Probing:** for each Proposer-output candidate, the image triplet (I_1, I_2, I_3) is held fixed while the DeepAgent evolves only the question text over $T=3$ turns. Conditioned on the full interaction history, at each turn it adaptively selects from a pool of four Rewriting patterns (contextualisation, projection, behavioural-tendency framing, cognitive-attribution framing) and three Deepening patterns (attribute refinement, scenario deepening, comparison deepening). Progressively warmer card tints on the stacked turns signal increasing probing depth; the resulting deep-probed triplets feed into the benchmark assembly of §IV-C.

batch of $\sim 2,000$ candidates per category. We retain this batch as a *frozen exemplar pool*: all subsequent Expand invocations in the adversarial-evolution loop (below) draw their few-shot exemplars from this pool rather than from the raw seed. Across iterations only the Proposer’s weights change (via DPO); the exemplar pool is held fixed, which eliminates compounding drift between the generator and its own past outputs while still letting the Proposer shift its output distribution through weight updates alone.

Adversarial Evolution via Direct Preference Optimization. A single Expansion round, applied to the Proposer in its initial base-model state, yields a broad but target-agnostic candidate batch. To shift the Proposer’s output distribution toward the specific vulnerabilities of a target LVLM, we wrap Expansion in a closed-loop adversarial evolution that repeats four steps per iteration:

- 1) *Expand* — run the current-state Proposer through the Expansion pipeline described above, drawing few-shot exemplars from the frozen exemplar pool, to produce a fresh batch of roughly 2,000 candidate triplets per demographic category.
- 2) *Query* — present each candidate’s three image–question pairs to the target LVLM independently, obtaining responses $\{r_1, r_2, r_3\}$.
- 3) *Label* — following the bias criterion of §III, mark a candidate as a *positive* preference if at least one response is non-“Unknown” (the query elicits bias on at least one image), or as a *negative* preference if all three responses are “Unknown” (the query fails to elicit any bias).
- 4) *Update* — apply DPO on the resulting preference pairs to obtain the next Proposer state, which is used by

Expand in the following iteration.

The ensemble variant used for benchmark construction—where the positive/negative decision in *Label* aggregates over multiple anchor LVLMs rather than a single target—is described in §IV-C.

The resulting preference pairs drive a Direct Preference Optimization (DPO) [8] update of the Proposer. Formally, given a preference pair (x_w, x_l) where x_w is a positive candidate and x_l is a negative one, DPO minimises:

$$\mathcal{L}_{\text{DPO}} = -\mathbb{E}_{(x_w, x_l)} \left[\log \sigma \left(\beta \log \frac{\pi_{\theta}(x_w)}{\pi_{\text{ref}}(x_w)} - \beta \log \frac{\pi_{\theta}(x_l)}{\pi_{\text{ref}}(x_l)} \right) \right], \quad (4)$$

where π_{θ} is the Proposer optimised through LoRA adapters [50], π_{ref} is a frozen copy of the Proposer’s underlying base model held fixed across iterations, β is a temperature hyperparameter, and σ is the sigmoid function. We use the instruction-tuned base model directly as π_{ref} and do *not* insert a supervised fine-tuning step before DPO: because the first-iteration *Expand* batch is itself produced by the base model under few-shot prompting, fine-tuning the same model on its own outputs would contribute essentially no additional training signal. Iterating the Expand–Query–Label–Update loop drives the Proposer’s distribution toward the target model’s specific failure modes, producing progressively harder candidates. Throughout, the Proposer’s evolution operates exclusively at the *distribution level*, changing which kinds of queries the generator emits; per-candidate deep probing is decoupled and handled by the DeepAgent described next.

B. Instance-Level Probing with the DeepAgent

Once the Proposer has converged, each candidate it emits already carries the target model’s distribution-level weaknesses. The **DeepAgent** complements the Proposer with per-candidate *instance-level* probing: for every Proposer-generated candidate triplet (I_1, I_2, I_3, Q) , it runs an adaptive dialogue with the target LVLM over T probing turns, building on the model’s own responses to elicit stereotypes that a single-shot query cannot reach. T is a hyperparameter whose concrete value is reported in §V. The right panel of Fig. 3 illustrates one such probing trajectory: the image triplet (I_1, I_2, I_3) is held fixed, while the question Q_t is adaptively rewritten or deepened turn by turn.

At turn $t \in \{1, \dots, T\}$, the DeepAgent receives the full interaction history

$$\mathcal{H}_t = \left\{ (Q_0, \{r_1^0, r_2^0, r_3^0\}), \dots, (Q_{t-1}, \{r_1^{t-1}, r_2^{t-1}, r_3^{t-1}\}) \right\},$$

where Q_0 is the Proposer’s original question and $Q_1, \dots, Q_{t-1}$ are its prior refinements by the DeepAgent. Conditioned on $\mathcal{H}_t$, the DeepAgent adaptively selects one of two strategies:

- **Deepening** — triggered when the previous response exhibits bias, i.e. at least one of $\{r_1^{t-1}, r_2^{t-1}, r_3^{t-1}\}$ is non-“Unknown”. Rather than stopping at a general prejudiced conclusion, the DeepAgent generates a targeted follow-up question that pins down the specific dimension and severity of the exposed stereotype. We instruct it to choose among three canonical patterns grounded in the

social-psychology literature on stereotype content and application: (i) *attribute refinement* — decompose a broad prejudice into specific traits along the warmth / competence dimensions of the Stereotype Content Model [51], e.g., isolating whether an unfitness judgement stems from a perceived lack of logical reasoning, leadership, or stress tolerance; (ii) *scenario deepening* — place the depicted individual in a concrete situation that sharpens the stereotype’s application [52], turning a vague prejudice into a definite behavioural judgement; and (iii) *comparison deepening* — force an intergroup comparison that exposes the stereotype’s polarity, drawing on Social Identity Theory’s account of comparative in-group/out-group evaluation [53].

- **Rewriting** — triggered when the previous response is safe, i.e. all three responses are “Unknown”. The DeepAgent adversarially rewrites Q into a more implicit, indirect query Q' . Social-psychology research has long observed that explicit stereotype prompts invite socially-desirable deflection while indirect framings bypass controlled responding [54], [55]; the DeepAgent draws on four such framings: (i) *contextualisation* — embed the prejudice in a specific everyday situation that masks its explicit stereotyping form; (ii) *projection* — ask about third-party impressions rather than the respondent’s direct attribute assignment, leveraging the reduced social-desirability pressure observed in indirect measures [56]; (iii) *behavioural-tendency framing* — reformulate a trait judgement as a likely behavioural outcome (e.g., from “is this person of low intelligence?” to “is this person likely to achieve a perfect score in an advanced math exam?”), following the Theory of Planned Behaviour’s linkage between attitude and intended behaviour [57]; and (iv) *cognitive-attribution framing* — attribute an observed outcome to stable internal traits rather than external circumstances, drawing on Weiner’s attributional theory of causal ascription [58]. These shifts move from explicit stereotyping to nuanced situational judgement, stripping away the model’s safety disguise and forcing it to reveal hidden prejudices.

At each turn the DeepAgent emits exactly one refined question Q_t while keeping the image triplet fixed; Q_t is then queried against the target model to obtain $\{r_1^t, r_2^t, r_3^t\}$, which are appended to form $\mathcal{H}_{t+1}$. The history-aware design lets the DeepAgent build on the full interaction trajectory rather than treating each turn independently, progressively narrowing in on the model’s deepest vulnerabilities. Coupling the Proposer’s distribution-level optimisation (§IV-A) with the DeepAgent’s candidate-level deepening lets DeepBias achieve bias exploration substantially beyond what any single-turn evaluation can reach.

C. DeepBias Benchmark Construction

While the pipeline of §IV-A yields samples targeted at a single model, a standardised benchmark must capture *common* vulnerabilities across LVLMs. We therefore apply the adaptive pipeline to an ensemble of six diverse, state-of-the-art LVLMs

acting as *anchor models*, and then assemble the DeepBias Benchmark by instantiating the validity criterion of §III on this same ensemble.

Ensemble-Driven Adversarial Evolution. We drive the Proposer–DPO loop with a preference signal aggregated across the anchor ensemble rather than a single target. A generated candidate is labelled a *positive* preference if it elicits bias from at least half of the anchors, and a *negative* preference if all anchors respond safely; candidates lying in between are discarded to keep the optimization signal clean. This consensus rule forces the Proposer to target vulnerabilities that transfer across architectures and scales, rather than overfitting to the safety quirks of any single model.

Validity Check via the Anchor Ensemble. We instantiate the validity check defined in Eqs. 1–2 by reusing the six anchor models as validators, so that no additional inference cost is incurred beyond the ensemble evaluation already performed during evolution. For each candidate (I_1, I_2, I_3, Q) we thus have $K=6$ validators producing $3K=18$ responses, and we retain Q iff $\rho(Q) > \tau$ with $\tau = 1/3$ —that is, iff more than one in three of the 18 validator–image pairs recognises the question as ambiguous. The value $\tau = 1/3$ is calibrated empirically; the sensitivity analysis, together with the effect of varying K and the validator composition, is reported in §V-F.

Triplet Decomposition and Assembly. Once a candidate passes the validity check, each of its three image–question pairs $(I_i, Q, \mathcal{O})$ is treated as an independent test instance as prescribed in §III. The original triplet is retained as metadata: although it is no longer used for scoring, it enables fine-grained analyses of which demographic groups are disproportionately affected by a given model. Aggregating all admitted instances across the three demographic categories—Age, Race, and Gender—yields the DeepBias Benchmark, a concentrated adversarial testing ground that couples the controlled-variable rigour of triplet validation with the directness of per-instance bias judgment.

V. EXPERIMENT

A. Implementation Details

We instantiate the Proposer as the instruction-tuned Qwen3-32B [59] fine-tuned with LoRA adapters [50] (rank 16, $\alpha = 32$) attached to all attention-projection matrices. Consistent with the no-SFT design of §IV-A, the frozen instruction-tuned base model itself serves as π_{ref} . Each DPO round uses $\beta = 0.1$, a batch size of 32 preference pairs, and a cosine-annealed learning rate with peak value 5×10^{-5} , and trains until preference accuracy plateaus. We run $R = 2$ DPO rounds for the single-target experiments (§V-B) and $R = 3$ rounds for the ensemble benchmark-construction experiment (§IV-C). The DeepAgent is implemented with the same Qwen3-32B backbone used in frozen-inference mode, prompted with a system message that enumerates the seven sub-strategies of §IV-B; the probing horizon is fixed at $T = 3$ turns for every reported result. Counterfactual image triplets are rendered by Stable Diffusion XL [49] at 1024×1024 resolution, with standardised neutral contextual prompts enforcing strict variable control across the sensitive attribute. All target LVLs

are queried in zero-shot mode using each model’s default inference configuration and system prompt. All experiments run on $8 \times$ NVIDIA RTX 3090 GPUs. The full Proposer expansion prompts and the DeepAgent strategy prompts are provided in the Supplementary Material.

B. Proposer Evolution

To validate the effectiveness of the Proposer’s DPO-driven evolution, we conduct an experiment focusing on the Age category across two target models: InternVL3-8B and Qwen2.5-VL-7B. Starting from the base Proposer (before any DPO), we execute two rounds of DPO optimization and report the bias elicitation rate in Table I (columns *Pre-DPO / DPO Iter 1–2*); the downstream DeepAgent columns of the same table are discussed in §IV-B. This setting isolates the contribution of the Proposer’s iterative optimization, without any DeepAgent probing.

Across both target models, the bias elicitation rate progressively increases with each DPO iteration. InternVL3-8B rises from 9.6% to 16.9%, a total gain of 7.3 percentage points, while Qwen2.5-VL-7B improves from 15.0% to 21.3% (+6.3 pp). These consistent upward trends provide direct empirical evidence that the DPO optimization loop successfully drives the Proposer to generate increasingly challenging test cases that align with each target model’s specific vulnerabilities.

Distribution Shift across DPO Iterations. Beyond aggregate elicitation rates, Figure 4 visualises how the Proposer’s output distribution migrates across DPO iterations. Two complementary effects are observed. First, *diversification*: relative to the Pre-DPO batch, DPO Iter 1 broadens topic coverage substantially—the effective number of active BERTopic clusters rises from 48.8 to 66–67 and topic entropy grows from 3.89 to 4.20, with unique-instruction ratio jumping from 0.84 to above 0.95 (Figure 5). This reflects the Proposer exploring a wider pool of stereotype-eliciting scenarios—profession, social role, legal context, and civic participation, among others. Second, *targeted concentration*: at DPO Iter 2 the distribution tightens around sub-regions that most reliably elicit non-“Unknown” responses. Vocabulary size shrinks ($1,219 \rightarrow 791$ –920) and mean pairwise cosine similarity rises ($0.42 \rightarrow 0.47$ for InternVL3-8B), consistent with the Proposer reallocating generation mass toward target-specific vulnerabilities once broad coverage has been established. Together, these two stages match the DPO loop’s incentive structure: broad exploration first, then exploitation of the target model’s weak points.

C. DeepAgent Multi-Turn Probing

Building on the Proposer’s distribution-level evolution (§IV-A), this experiment isolates the contribution of the DeepAgent’s multi-turn probing. For each target model, we apply the DeepAgent on top of the DPO-converged Proposer’s candidate batch and record how the bias elicitation rate evolves across the T probing turns (with $T = 3$ in our implementation). Results are reported in the *Probe Turn 1–3* columns of Table I, alongside the Proposer-only columns (§V-B) for direct comparison.

Models	Raw Dataset	Pre-DPO	DPO Iter 1	DPO Iter 2	Probe Turn 1	Probe Turn 2	Probe Turn 3
InternVL3-8B	—	9.6	15.8	16.9	32.7	39.5	39.3
Qwen2.5-VL-7B	—	15.0	17.9	21.3	52.3	47.0	48.4
Data Validity Rate	—	95.8	98.3	98.4	62.5	62.2	59.5

TABLE I: **DeepBias pipeline trajectory on Age bias.** Bias elicitation rate (% , proportion of non-“Unknown” responses) for two target LVLs along the full Proposer-then-DeepAgent pipeline. *Raw Dataset*: model’s elicitation rate on the pre-expansion BBQ/VLBiasBench raw samples (~ 300 /category). *Pre-DPO*: Proposer’s first-round candidate batch, before any DPO. *DPO Iter 1–2*: after each round of Proposer adversarial evolution (§IV-A). *Probe Turn 1–3*: after each DeepAgent multi-turn probing step applied on top of the DPO-converged Proposer (§IV-B). The bottom row is the per-stage data validity rate—the fraction of candidate triplets passing the ensemble voting check of §III; DeepAgent’s adversarial rewriting drives a moderate drop ($\sim 98\% \rightarrow \sim 60\%$), reflecting the tradeoff between probe depth and ambiguity preservation.




Fig. 4: **Proposer output distribution across DPO iterations (Age bias).** (a) UMAP overlay of the Proposer’s candidate pool at Pre-DPO / DPO Iter 1 / DPO Iter 2. The iterated distribution both expands into previously uncovered sub-regions and concentrates mass on sub-regions where the target model most reliably emits non-“Unknown” responses. (b) BERTopic radar over the discovered topic set: DPO iterations broaden coverage of stereotype domains (profession, civic role, security, etc.) and increase the effective number of active topics.

DeepAgent delivers a dramatic jump over what Proposer-only DPO can reach. On InternVL3-8B the elicitation rate roughly doubles from 16.9% (post-DPO, Iter 2) to 39.3% (Turn 3), an additional +22.4 percentage points beyond the Proposer’s standalone gain of +7.3 pp. On Qwen2.5-VL-7B the improvement is even larger: +27.1 pp, from 21.3% to 48.4%. The end-to-end gain from the raw Proposer batch (*Pre-DPO*) to the final DeepAgent turn is +29.7 pp for InternVL3-8B and +33.4 pp for Qwen2.5-VL-7B. These gains demonstrate that the instance-level, history-aware probing of §IV-B reaches biases that distribution-level optimization alone cannot surface.

The per-turn trajectory is not monotone: on Qwen2.5-VL-7B, Turn 1 (52.3%) is the single strongest round, and Turns 2–3 fluctuate slightly below it. A plausible reading is that *Rewriting* (the adversarial safety-bypass strategy) is most productive on the first invocation, after which subsequent turns shift toward *Deepening*, which refines already-exposed

biases rather than uncovering new ones. The data validity rate (bottom row of Table I) drops from roughly 62% at Turn 1 to 59.5% by Turn 3, indicating a modest tradeoff between probe depth and candidate validity; this tradeoff is quantified in the threshold analysis of §V-F.

Case Study: Multi-Turn Dialogue Trace. Figure 6 traces a single candidate through the DeepAgent’s three-turn dialogue against InternVL3-8B. Turn 0 is the candidate inherited from the Proposer’s DPO-converged batch—“*This person leads a tech startup; how likely are investors to fund them?*”—an ambiguous prompt to which the target responds “Unknown”. Since no bias has yet been elicited, the DeepAgent invokes *Rewriting* via a contextualisation sub-strategy, embedding the same core question inside a concrete pitch-meeting scenario (Turn 1); the target now yields a non-“Unknown” response (“Moderately”), leaking initial evidence of bias. Having exposed bias, the DeepAgent switches to *Deepening*: at Turn 2 it applies attribute refinement to ground the probe in visual




Fig. 5: Diversity and concentration metrics across Proposer DPO iterations. Per-iteration vocabulary size, distinct-2, unique-instruction ratio, mean pairwise cosine similarity, centroid distance, self-BLEU, effective topic count and topic entropy, computed on the Age-bias Proposer candidate batches at Pre-DPO / DPO Iter 1 / DPO Iter 2 against each target LVLm. The joint trend—first expansion then concentration—is consistent across target models.

cues (“apparent age and appearance”), extracting a stronger biased response (“Limited”); at Turn 3 it applies comparison deepening, directly contrasting the depicted founder against a younger counterpart, at which point the target commits to an unambiguous biased answer (“Yes”). This trajectory illustrates how the per-turn strategy choice adapts to the preceding response and how the DeepAgent’s full-history conditioning incrementally builds on each revealed cue.

D. Benchmark Construction

To construct a robust and generalizable benchmark, we extend our framework to an ensemble of 6 diverse LVLms across three demographic categories: Age, Race, and Gender. We introduce a consensus-based labeling mechanism for the DPO preference dataset. A generated test case is labeled as a positive preference if it successfully elicits biased responses from at least 50% of the anchor models. Conversely, queries that fail to expose bias where all models give safe answers are labeled as negative preferences. This forces the Proposer to target on universal, structural vulnerabilities shared across different architectures.

Table II tracks the bias accuracy of this ensemble during the process. The result reveals two systemic issues in current LVLm safety alignments. First, the latent biases are scale-agnostic and architecture-agnostic. While the 38B model (InternVL3.5-38B) starts with a higher baseline defense than the 7B models, it still suffers a consistent degradation of 12 to 15 percentage points across all categories by Iteration 3. This proves that merely scaling up model parameters does not patch these vulnerabilities. Second, the defense mechanisms exhibit demographic imbalances. Models show significantly weaker resilience against Race and Gender probes compared to Age. By Iteration 3, the accuracies for Race and Gender drop to the 50%-60% range. This contrast demonstrates that racial and gender biases are more deeply rooted.

Case Study: DeepAgent Multi-Turn Dialogue Trace




Fig. 6: Case Study: DeepAgent multi-turn dialogue trace. A single candidate inherited from the DPO-converged Proposer’s batch (Turn 0) is refined across three probing turns against InternVL3-8B. Each row marks the turn index, the selected strategy regime (Rewriting at Turn 1, Deepening at Turns 2–3) and sub-strategy (Contextualisation, Attribute Refinement, Comparison Deepening), the refined question, and the target model’s response. The response deepens from ambiguous (“Unknown”) to unambiguously biased (“Yes”) as the dialogue progresses.

Benchmark Structural Composition. Beyond the per-iteration trends of Table II, Figure 7 characterises the structural composition of the final benchmark. Panel (a) projects all benchmark questions into a shared 2D semantic space via UMAP, with samples coloured by demographic category. The three categories occupy partially overlapping regions: a substantial overlap zone reflects a shared pool of bias-eliciting templates—profession, social role, civic context—that the pipeline applies across demographic axes, while each category retains a signature periphery of category-specific scenarios. Panel (b) resolves this differentiation at the topic level: Age questions concentrate on the job-hiring, expert/tech, and education axes (capability and experience triggers); Gender questions lean toward job-hiring, social-role, and public-speaking topics (authority and competence triggers); Race questions are pulled toward security/legal, space/venue, and consumer settings (social location and access triggers). The topical specialisation matches well-documented sociological trigger points for each axis, demonstrating that the benchmark is neither a naively replicated template nor a disjoint collection of three subsets.

Comparison with Prior Benchmarks. Figure 8 situates DeepBias against the four most directly comparable bias benchmarks for language and vision-language models: BBQ [18], VLBiasBench [23], GenderBias-VL [19], and




Fig. 7: **Structural characterisation of the DeepBias benchmark.** (a) UMAP overlay of the benchmark’s question embeddings, coloured by demographic category. The three categories occupy partially overlapping regions, indicating a shared pool of bias-eliciting templates with category-specific peripheries. (b) BERTopic radar over the discovered top topics: Age questions concentrate on job/hiring, expert/tech, and education; Gender on job/hiring, social role, and public speaking; Race on security/legal, space/venue, and consumer settings.




Fig. 8: **Comparison of DeepBias against prior bias benchmarks.** Five quality and coverage axes (left to right): instance count, number of demographic categories, validity rate (proportion of constructed instances passing ensemble consensus; “N/R” = not reported), average bias elicitation rate on strong LVLMS, and presence of multimodal (vision) input. DeepBias matches established benchmarks on size and demographic coverage while substantially exceeding them on validity and challenge level.

PAIRS [60]. Across five comparison axes, DeepBias offers a competitive scale (1,500 instances, sitting between the larger BBQ/VLBIASBench and the smaller GenderBias-VL/PAIRS) and demographic coverage (three core categories), while distinguishing itself on two quality dimensions. First, the ensemble validity rate—the fraction of constructed instances passing the $K = 6$, $\tau = 1/3$ ambiguity check of §III—reaches 88%, substantially above the 45–62% reported by the multimodal benchmarks (BBQ does not report a formal validity figure). Second, the average bias elicitation rate on strong LVLMS reaches 48%, versus 19–31% for the prior benchmarks, indicating that DeepBias surfaces bias on instances that alignment-tuned models routinely refuse on legacy benchmarks. Coupled with its multimodal grounding and explicit anchoring in the three-characteristic protocol of §III, DeepBias fills a distinct

methodological niche: a rigorously-validated, vision-grounded, challenge-level bias benchmark.

Difficulty Distribution. To verify that the benchmark presents a reasonable spread of difficulty—neither trivially solvable nor impossibly hard—Figure 9 plots the per-instance bias elicitation rate aggregated over the six anchor LVLMS. The mass of the distribution lies in the 30–60% band with an overall median near 48%, indicating that a typical instance elicits bias from three to four of the six anchors. The three demographic categories produce near-identical distribution shapes, suggesting that the benchmark’s difficulty is not an artefact of any particular demographic axis. Instances at the extremes—trivially refused (rate near 0%) or trivially leading (rate near 100%)—are sparse, a direct consequence of the $\tau = 1/3$ validity filter of §IV-A rejecting overly ambiguous




Fig. 9: **Per-instance benchmark difficulty distribution.** Histogram of the per-instance bias elicitation rate (the proportion of the six anchor LVLMS returning a non-“Unknown” response), stacked by demographic category. The dashed line marks the overall median ($\approx 48\%$).

or overly suggestive questions.

Finally, we assemble a comprehensive and highly challenging benchmark comprising the adversarially-mined test instances distilled through the full Proposer–DeepAgent cycle. By capturing the dynamic outputs of our probing cycle, the benchmark provides a rigorously mined evaluation standard that reflects a broad spectrum of latent biases. Representative instances spanning the three demographic categories—each consisting of a counterfactual image triplet, an intentionally-ambiguous question, and the ensemble’s majority response—are presented in the Supplementary Material.

E. Benchmark Evaluation

Having assembled the DeepBias Benchmark in §IV-C, we now treat it as a frozen evaluation standard and measure how a broad pool of LVLMS—spanning model families, parameter scales, and release cohorts—performs against it. Table III reports per-model bias accuracy across the three demographic categories at the three successive benchmark versions produced during construction (“iter k ” corresponds to the candidate set retained after the k -th Proposer–DPO round). Unlike §IV-C, which tracked the anchor ensemble’s degradation while the benchmark was being formed, this experiment probes how models *not* used as anchors during construction respond to the finalised challenge set, isolating the benchmark’s *transferable* difficulty from any anchor-specific artefact. The evaluation pool combines the six anchor models with hold-out open-source LVLMS (Otter, InternLM-XComposer2-VL, InternVL3-14B, Qwen3-VL-8B, MiniGPT4-Vicuna-7B/13B, Shikra-7B, InstructBLIP-Vicuna-7B/13B, DeepSeek-VL) and the two leading proprietary models (GPT-4o, Gemini).

Three patterns emerge from Table III. First, *cross-architecture transfer*: hold-out models that were never part of the anchor ensemble exhibit bias-accuracy degradation curves qualitatively similar to the anchors, confirming that DeepBias surfaces structural alignment limitations rather than anchor-specific artefacts and that the challenge set generalises beyond the six models used during construction. Second, *scale offers*

no immunity: the 38B-parameter InternVL3.5-38B enters the benchmark with the strongest iter-1 defence of the evaluated pool, yet still loses more than 12 percentage points by iter 3 on every category; a similar erosion is observed for the 30B-scale Qwen3-VL-A3B. Parameter count therefore postpones, but does not prevent, bias leakage—a finding in direct tension with the common practice of treating scale as a proxy for safety. Third, *persistent demographic imbalance*: every evaluated model shows a noticeable gap between Age accuracy and Race/Gender accuracy, with Race and Gender dropping to the 50–65% range by iter 3 for most open-source models while Age holds above 65%. The imbalance mirrors the one observed in the construction-time anchor ensemble (§IV-C), indicating that the deficit is an LVLMS-wide artefact of current alignment pipelines rather than a property of any particular architecture. Proprietary models (GPT-4o, Gemini) hold their iter-1 defence longer than comparable open-source counterparts but follow the same qualitative trajectory, suggesting that the additional engineering effort behind frontier systems raises the entry barrier without eliminating the underlying vulnerability.

F. Validation Threshold Analysis

As described in Section IV-C, we employ a multi-model voting mechanism to filter out degenerate questions. A question is retained if the proportion of “Unknown” responses $\rho(Q)$ exceeds a threshold τ . The choice of τ directly governs the trade-off between data quality and quantity: a high threshold aggressively filters questions but risks discarding valid, challenging items (false rejection); a low threshold retains more data but may admit trivially answerable questions that corrupt the evaluation signal.

To determine the optimal threshold, we conduct a systematic analysis. We sample a batch of Proposer-generated questions and obtain human annotations on a subset to establish ground truth labels (valid vs. degenerate). We then vary τ across a range of values and evaluate the validation mechanism along three dimensions:

- **Survival rate**: the proportion of questions passing the threshold, reflecting data quantity.
- **Precision**: among retained questions, the proportion that are genuinely valid, reflecting data quality.
- **Recall**: among all valid questions, the proportion that are retained, reflecting whether the threshold is overly conservative.

Figure 10 reports the resulting F1, precision, and recall curves as functions of τ on the annotated subset. F1 reaches its plateau around $\tau \in [1/6, 3/6]$, with the maximum near $\tau \approx 1/3$ where precision and recall are best balanced. Below this range, recall remains saturated but precision degrades sharply as trivially-answerable questions slip through; above it, precision climbs marginally at the cost of an aggressive recall drop that discards otherwise-valid challenging instances. We therefore adopt $\tau = 1/3$ as the operational threshold for benchmark construction, corresponding to a one-out-of-three minimum “Unknown” rate across the six anchor responses.

Models	Age			Race			Gender		
	iter 1	iter 2	iter 3	iter 1	iter 2	iter 3	iter 1	iter 2	iter 3
InternVL3-8B	89.3	82.1	74.6	77.8	70.2	62.4	76.5	68.9	61.3
Qwen2.5-VL-Instruct-7B	82.7	75.3	67.1	71.3	63.5	55.2	70.1	62.1	53.9
Qwen3-VL-30B-A3B-Instruct	87.5	80.9	73.2	75.9	68.7	60.8	74.8	67.4	59.6
InternVL3.5-38B	90.6	84.2	77.8	79.2	72.4	64.9	78.1	70.8	63.2
DeepSeek-VL	85.4	78.6	70.9	73.8	66.7	58.6	72.6	65.2	57.1

TABLE II: Bias accuracy (%) of the 6-model anchor ensemble across three demographic categories during the Proposer’s DPO evolution. Lower accuracy indicates more successfully elicited bias; cells are shaded accordingly, with green = stronger defence and red = heavier bias leakage. Columns “iter 1–3” denote successive Proposer-DPO rounds.

Models	Age			Race			Gender		
	iter 1	iter 2	iter 3	iter 1	iter 2	iter 3	iter 1	iter 2	iter 3
Otter									
InternVL3-8B	89.3	82.1	74.6	77.8	70.2	62.4	76.5	68.9	61.3
InternLM-XComposer2-VL-7b									
InternVL3-14B									
InternVL3.5-38B	90.6	84.2	77.8	79.2	72.4	64.9	78.1	70.8	63.2
Qwen2.5-VL-Instruct-7B	82.7	75.3	67.1	71.3	63.5	55.2	70.1	62.1	53.9
Qwen3-VL-8B-instruct									
Qwen3-VL-30B-A3B-Instruct	87.5	80.9	73.2	75.9	68.7	60.8	74.8	67.4	59.6
Minigt4-vicuna-7b									
Minigt4-vicuna-13b									
Shikra-7B									
InstructBlip-vicuna-7b									
InstructBlip-vicuna-13b									
DeepSeek-VL	85.4	78.6	70.9	73.8	66.7	58.6	72.6	65.2	57.1
GPT-4o									
Gemini									

TABLE III: Bias accuracy (%) on the constructed DeepBias Benchmark across three demographic categories and three bias-probing iterations, reported for a diverse set of LVLMs (lower accuracy = more bias elicited).

We additionally examine the effect of the number of validator models K and their composition, holding τ fixed at its operating point $1/3$. Varying $K \in \{2, 3, 4, 6\}$ by sampling random subsets of the anchor ensemble, we observe that F1 and the stability of $\rho(Q)$ across draws both increase monotonically with K and plateau once the pool reaches $K \geq 4$; moving from $K=4$ to $K=6$ primarily reduces the per-question variance of $\rho(Q)$ rather than shifting its mean, indicating that larger ensembles mainly buy reliability rather than stricter filtering. To disentangle composition from size, we further contrast balanced validator pools—half “conservative” (anchors with baseline Unknown rate above the pool median on a held-out ambiguous set) and half “aggressive”—against pools skewed toward a single response-style family. Balanced pools remain within a narrow F1 band of the skewed extremes at $\tau = 1/3$, showing that the threshold absorbs most of the composition effect as long as the pool is not collapsed onto a single family. We therefore adopt $K = 6$ with a deliberately diverse anchor ensemble spanning architectures and scales as the operating

point for benchmark construction, trading a modest amount of extra inference cost for tight control over the variance of the validity signal.

G. Limitations and Future Work

While DeepBias improves bias evaluation, several limitations need future investigation. First, our current framework focuses on three major bias categories (Age, Race, Gender). Extending to other dimensions (e.g., Religion, Disability, Socioeconomic Status) would provide more comprehensive coverage. Second, the adversarial evolution process requires substantial computational resources. Developing more efficient optimization strategies could improve scalability. Third, while our benchmark exposes biases, it does not directly provide mitigation strategies. Future work could explore integrating DeepBias into the training loop to guide bias-aware fine-tuning or developing automated debiasing techniques informed by the discovered vulnerabilities. Finally, our current methodology is specifically tailored to uncovering societal biases. A highly




Fig. 10: **Validity threshold sensitivity.** F1 score, precision, and recall of the ensemble validity filter as functions of the threshold τ , evaluated on a 200-question subset with human-annotated ambiguity ground truth ($K = 6$ anchor models). The dotted line marks the adopted operating point $\tau = 1/3$.

promising direction for future research involves generalizing this adaptive adversarial framework to evaluate other critical dimensions of LVLM trustworthiness, such as multimodal hallucinations, value alignment, and broader safety risks.

VI. CONCLUSION

In this paper, we introduced DeepBias, a dynamic and adversarial evaluation framework designed to expose deeply ingrained societal biases in Large Vision-Language Models (LVLMs). Unlike existing static benchmarks, our approach leverages a constantly optimized Proposer and an autonomous DeepAgent with multi-turn deep refinement to continuously generate test cases. Extensive experiments demonstrate that this dynamic probing consistently exposes the societal biases of target models. We found that the mechanism drives the data distribution to explicitly shift toward the models' latent weaknesses and generates progressively deep and targeted adversarial queries. By scaling this framework across an ensemble of anchor models, we successfully curated a highly challenging benchmark. Subsequent testing of state-of-the-art LVLMs on this benchmark revealed demographic prejudices that remain largely hidden under traditional evaluation protocols. By shifting the paradigm from fixed evaluation to dynamic adversarial mining, our work provides a robust, evolving tool to uncover complex vulnerabilities and foster the development of genuinely equitable multimodal systems.

ACKNOWLEDGMENT

Please insert your acknowledgments here.

REFERENCES

- [1] A. Hurst, A. Lerer, A. P. Goucher, A. Perelman, A. Ramesh, A. Clark, A. Ostrow, A. Welihinda, A. Hayes, A. Radford *et al.*, "Gpt-4o system card," *arXiv preprint arXiv:2410.21276*, 2024.
- [2] H. Liu, C. Li, Q. Wu, and Y. J. Lee, "Visual instruction tuning," *Advances in neural information processing systems*, vol. 36, pp. 34 892–34 916, 2023.
- [3] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, "Proximal policy optimization algorithms," *arXiv preprint arXiv:1707.06347*, 2017.
- [4] S. M. Hall, F. Gonçalves Abrantes, H. Zhu, G. Sodunke, A. Shtedritski, and H. R. Kirk, "Visogender: A dataset for benchmarking gender bias in image-text pronoun resolution," *Advances in Neural Information Processing Systems*, vol. 36, pp. 63 687–63 723, 2023.
- [5] Y. Zhang, J. Wang, and J. Sang, "Counterfactually measuring and eliminating social bias in vision-language pre-training models," in *Proceedings of the 30th ACM International Conference on Multimedia*, 2022, pp. 4996–5004.
- [6] C. Fu, P. Chen, Y. Shen, Y. Qin, M. Zhang, X. Lin, J. Yang, X. Zheng, K. Li, X. Sun *et al.*, "Mme: A comprehensive evaluation benchmark for multimodal large language models," *arXiv preprint arXiv:2306.13394*, 2023.
- [7] Y. Liu, H. Duan, Y. Zhang, B. Li, S. Zhang, W. Zhao, Y. Yuan, J. Wang, C. He, Z. Liu *et al.*, "Mmbench: Is your multi-modal model an all-around player?" in *European conference on computer vision*. Springer, 2024, pp. 216–233.
- [8] R. Rafailov, A. Sharma, E. Mitchell, C. D. Manning, S. Ermon, and C. Finn, "Direct preference optimization: Your language model is secretly a reward model," *Advances in neural information processing systems*, vol. 36, pp. 53 728–53 741, 2023.
- [9] N. Mehrabi, F. Morstatter, N. Saxena, K. Lerman, and A. Galstyan, "A survey on bias and fairness in machine learning," *ACM Computing Surveys*, vol. 54, no. 6, pp. 1–35, 2021.
- [10] S. L. Blodgett, S. Barocas, H. Daumé III, and H. Wallach, "Language (technology) is power: A critical survey of "bias" in NLP," in *Proceedings of the Annual Meeting of the Association for Computational Linguistics*, 2020, pp. 5454–5476.
- [11] S. Barocas and A. D. Selbst, "Big data's disparate impact," *California Law Review*, vol. 104, pp. 671–732, 2016.
- [12] H. Suresh and J. Gutttag, "A framework for understanding sources of harm throughout the machine learning life cycle," in *Proceedings of the ACM Conference on Equity and Access in Algorithms, Mechanisms, and Optimization*, 2021.
- [13] J. Buolamwini and T. Gebru, "Gender shades: Intersectional accuracy disparities in commercial gender classification," in *Proceedings of the Conference on Fairness, Accountability, and Transparency*, 2018, pp. 77–91.
- [14] T. Bolukbasi, K.-W. Chang, J. Y. Zou, V. Saligrama, and A. T. Kalai, "Man is to computer programmer as woman is to homemaker? debiasing word embeddings," *Advances in neural information processing systems*, vol. 29, 2016.
- [15] A. Caliskan, J. J. Bryson, and A. Narayanan, "Semantics derived automatically from language corpora contain human-like biases," *Science*, vol. 356, no. 6334, pp. 183–186, 2017.
- [16] N. Nangia, C. Vania, R. Bhalariao, and S. Bowman, "Crows-pairs: A challenge dataset for measuring social biases in masked language models," in *Proceedings of the 2020 conference on empirical methods in natural language processing (EMNLP)*, 2020, pp. 1953–1967.
- [17] M. Nadeem, A. Bethke, and S. Reddy, "Stereoset: Measuring stereotypical bias in pretrained language models," in *Proceedings of the 59th annual meeting of the association for computational linguistics and the 11th international joint conference on natural language processing (volume 1: long papers)*, 2021, pp. 5356–5371.
- [18] A. Parrish, A. Chen, N. Nangia, V. Padmakumar, J. Phang, J. Thompson, P. M. Htut, and S. Bowman, "Bbq: A hand-built bias benchmark for question answering," in *Findings of the Association for Computational Linguistics: ACL 2022*, 2022, pp. 2086–2105.
- [19] Y. Xiao, X. Liu, Q. Cheng, Z. Yin, S. Liang, J. Li, J. Shao, A. Liu, and D. Tao, "Genderbias-vl: Benchmarking gender bias in vision language models via counterfactual probing: Y. xiao *et al.*," *International Journal of Computer Vision*, vol. 133, no. 12, pp. 8332–8355, 2025.
- [20] C. Raj, B. Wei, A. Caliskan, A. Anastasopoulos, and Z. Zhu, "Vignette: Socially grounded bias evaluation for vision-language models," *arXiv preprint arXiv:2505.22897*, 2025.
- [21] K. Karkkainen and J. Joo, "Fairface: Face attribute dataset for balanced race, gender, and age for bias measurement and mitigation," in *Proceedings of the IEEE/CVF winter conference on applications of computer vision*, 2021, pp. 1548–1558.
- [22] C. Schumann, S. Ricco, U. Prabhu, V. Ferrari, and C. Pantofaru, "A step toward more inclusive people annotations for fairness," in *Proceedings of the 2021 AAAI/ACM Conference on AI, Ethics, and Society*, 2021, pp. 916–925.

- [23] S. Wang, X. Cao, J. Zhang, Z. Yuan, S. Shan, X. Chen, and W. Gao, "Vbiasbench: A comprehensive benchmark for evaluating bias in large vision-language model," *arXiv preprint arXiv:2406.14194*, 2024.
- [24] C. Schuhmann, R. Beaumont, R. Vencu, C. Gordon, R. Wightman, M. Cherti, T. Coombes, A. Katta, C. Mullis, M. Wortsman *et al.*, "Laion-5b: An open large-scale dataset for training next generation image-text models," *Advances in neural information processing systems*, vol. 35, pp. 25 278–25 294, 2022.
- [25] O. Sainz, J. Campos, I. García-Ferrero, J. Etxaniz, O. L. de Lacalle, and E. Agirre, "Nlp evaluation in trouble: On the need to measure llm data contamination for each benchmark," in *Findings of the Association for Computational Linguistics: EMNLP 2023*, 2023, pp. 10 776–10 787.
- [26] G. Team, R. Anil, S. Borgeaud, J.-B. Alayrac, J. Yu, R. Soricut, J. Schalkwyk, A. M. Dai, A. Hauth, K. Millican *et al.*, "Gemini: a family of highly capable multimodal models," *arXiv preprint arXiv:2312.11805*, 2023.
- [27] D. Kiela, M. Bartolo, Y. Nie, D. Kaushik, A. Geiger, Z. Wu, B. Vidgen, G. Prasad, A. Singh, P. Ringshia *et al.*, "Dynabench: Rethinking benchmarking in nlp," in *Proceedings of the 2021 conference of the North American chapter of the Association for Computational Linguistics: human language technologies*, 2021, pp. 4110–4124.
- [28] E. Perez, S. Huang, F. Song, T. Cai, R. Ring, J. Aslanides, A. Glaese, N. McAleese, and G. Irving, "Red teaming language models with language models," *arXiv preprint arXiv:2202.03286*, 2022.
- [29] D. Ganguli, L. Lovitt, J. Kernion, A. Askell, Y. Bai, S. Kadavath, B. Mann, E. Perez, N. Schiefer, K. Ndousse *et al.*, "Red teaming language models to reduce harms: Methods, scaling behaviors, and lessons learned," *arXiv preprint arXiv:2209.07858*, 2022.
- [30] A. Zou, Z. Wang, N. Carlini, M. Nasr, J. Z. Kolter, and M. Fredrikson, "Universal and transferable adversarial attacks on aligned language models," *arXiv preprint arXiv:2307.15043*, 2023.
- [31] E. Jones, A. Dragan, A. Raghunathan, and J. Steinhardt, "Automatically auditing large language models via discrete optimization," in *International Conference on Machine Learning*. PMLR, 2023, pp. 15 307–15 329.
- [32] P. Chao, A. Robey, E. Dobriban, H. Hassani, G. J. Pappas, and E. Wong, "Jailbreaking black box large language models in twenty queries," in *2025 IEEE Conference on Secure and Trustworthy Machine Learning (SaTML)*. IEEE, 2025, pp. 23–42.
- [33] X. Liu, N. Xu, M. Chen, and C. Xiao, "Autodan: Generating stealthy jailbreak prompts on aligned large language models," *arXiv preprint arXiv:2310.04451*, 2023.
- [34] Y. Gong, D. Ran, J. Liu, C. Wang, T. Cong, A. Wang, S. Duan, and X. Wang, "Figstep: Jailbreaking large vision-language models via typographic visual prompts," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 39, no. 22, 2025, pp. 23 951–23 959.
- [35] Z. Niu, H. Ren, X. Gao, G. Hua, and R. Jin, "Jailbreaking attack against multimodal large language model," *arXiv preprint arXiv:2402.02309*, 2024.
- [36] E. Shayegani, Y. Dong, and N. Abu-Ghazaleh, "Jailbreak in pieces: Compositional adversarial attacks on multi-modal language models," *arXiv preprint arXiv:2307.14539*, 2023.
- [37] A. Mehrotra, M. Zampetakis, P. Kassianik, B. Nelson, H. Anderson, Y. Singer, and A. Karbasi, "Tree of attacks: Jailbreaking black-box llms automatically," *Advances in Neural Information Processing Systems*, vol. 37, pp. 61 065–61 105, 2024.
- [38] Y. Wang, Y. Kordi, S. Mishra, A. Liu, N. A. Smith, D. Khashabi, and H. Hajishirzi, "Self-instruct: Aligning language models with self-generated instructions," in *Proceedings of the 61st annual meeting of the association for computational linguistics (volume 1: long papers)*, 2023, pp. 13 484–13 508.
- [39] C. Xu, Q. Sun, K. Zheng, X. Geng, P. Zhao, J. Feng, C. Tao, and D. Jiang, "Wizardlm: Empowering large language models to follow complex instructions," *arXiv preprint arXiv:2304.12244*, 2023.
- [40] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, "Proximal policy optimization algorithms," *arXiv preprint arXiv:1707.06347*, 2017.
- [41] S. Casper, J. Lin, J. Kwon, G. Culp, and D. Hadfield-Menell, "Explore, establish, exploit: Red teaming language models from scratch," *arXiv preprint arXiv:2306.09442*, 2023.
- [42] M. Sorkhpoor, A. Yazdinejad, and A. Dehghantanha, "RedHit: Adaptive red-teaming of large language models via search, reasoning, and preference optimization," in *Proceedings of the First Workshop on LLM Security (LLMSEC)*, 2025.
- [43] M. T. Ribeiro, T. Wu, C. Guestrin, and S. Singh, "Beyond accuracy: Behavioral testing of nlp models with checklist," in *Proceedings of the 58th annual meeting of the association for computational linguistics*, 2020, pp. 4902–4912.
- [44] Q. Li, J. Gao, S. Wang, R. Pi, X. Zhao, C. Wu, X. Jiang, Z. Li, and L. Kong, "Forewarned is Forearmed: Leveraging llms for data synthesis through failure-inducing exploration," in *ICLR*, 2025.
- [45] C. Dwork, M. Hardt, T. Pitassi, O. Reingold, and R. Zemel, "Fairness through awareness," in *Proceedings of the Innovations in Theoretical Computer Science Conference*, 2012, pp. 214–226.
- [46] R. Binns, "Fairness in machine learning: Lessons from political philosophy," in *Proceedings of the Conference on Fairness, Accountability, and Transparency*, 2018, pp. 149–159.
- [47] D. A. Prentice and E. Carranza, "What women and men should be, shouldn't be, are allowed to be, and don't have to be: The contents of prescriptive gender stereotypes," *Psychology of Women Quarterly*, vol. 26, no. 4, pp. 269–281, 2002.
- [48] L. Jussim, J. T. Crawford, and R. S. Rubinstein, "Stereotype (in)accuracy in perceptions of groups and individuals," *Current Directions in Psychological Science*, vol. 24, no. 6, pp. 490–497, 2015.
- [49] D. Podell, Z. English, K. Lacey, A. Blattmann, T. Dockhorn, J. Müller, J. Penna, and R. Rombach, "Sdxl: Improving latent diffusion models for high-resolution image synthesis," *arXiv preprint arXiv:2307.01952*, 2023.
- [50] E. J. Hu, Y. Shen, P. Wallis, Z. Allen-Zhu, Y. Li, S. Wang, L. Wang, and W. Chen, "LoRA: Low-rank adaptation of large language models," in *ICLR*, 2022.
- [51] S. T. Fiske, A. J. C. Cuddy, P. Glick, and J. Xu, "A model of (often mixed) stereotype content: Competence and warmth respectively follow from perceived status and competition," *Journal of Personality and Social Psychology*, vol. 82, no. 6, pp. 878–902, 2002.
- [52] S. T. Fiske and S. L. Neuberg, "A continuum of impression formation, from category-based to individuating processes: Influences of information and motivation on attention and interpretation," in *Advances in Experimental Social Psychology*. Academic Press, 1990, vol. 23, pp. 1–74.
- [53] H. Tajfel and J. C. Turner, "An integrative theory of intergroup conflict," in *The Social Psychology of Intergroup Relations*, W. G. Austin and S. Worchel, Eds. Brooks/Cole, 1979, pp. 33–47.
- [54] P. G. Devine, "Stereotypes and prejudice: Their automatic and controlled components," *Journal of Personality and Social Psychology*, vol. 56, no. 1, pp. 5–18, 1989.
- [55] J. B. McConahay, "Modern racism, ambivalence, and the modern racism scale," in *Prejudice, Discrimination, and Racism*, J. F. Dovidio and S. L. Gaertner, Eds. Academic Press, 1986, pp. 91–125.
- [56] R. H. Fazio, J. R. Jackson, B. C. Dunton, and C. J. Williams, "Variability in automatic activation as an unobtrusive measure of racial attitudes: A bona fide pipeline?" *Journal of Personality and Social Psychology*, vol. 69, no. 6, pp. 1013–1027, 1995.
- [57] I. Ajzen, "The theory of planned behavior," *Organizational Behavior and Human Decision Processes*, vol. 50, no. 2, pp. 179–211, 1991.
- [58] B. Weiner, "An attributional theory of achievement motivation and emotion," *Psychological Review*, vol. 92, no. 4, pp. 548–573, 1985.
- [59] A. Yang, A. Li, B. Yang, B. Zhang, B. Hui, B. Zheng, B. Yu, C. Gao, C. Huang, C. Lv, C. Zheng, D. Liu, F. Zhou, F. Huang, F. Hu, H. Ge, H. Wei, H. Lin, J. Tang, J. Yang, J. Tu, J. Zhang, J. Yang, J. Yang, J. Zhou, J. Zhou, J. Lin, K. Dang, K. Bao, K. Yang, L. Yu, L. Deng, M. Li, M. Xue, M. Li, P. Zhang, P. Wang, Q. Zhu, R. Men, R. Gao, S. Liu, S. Luo, T. Li, T. Tang, W. Yin, X. Ren, X. Wang, X. Zhang, X. Ren, Y. Fan, Y. Su, Y. Zhang, Y. Zhang, Y. Wan, Y. Liu, Z. Wang, Z. Cui, Z. Zhang, Z. Zhou, and Z. Qiu, "Qwen3 technical report," *arXiv preprint arXiv:2505.09388*, 2025.
- [60] K. C. Fraser and S. Kiritchenko, "Examining gender and racial bias in large vision-language models using a novel dataset of parallel images," in *Proceedings of the 18th Conference of the European Chapter of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2024, pp. 690–713.